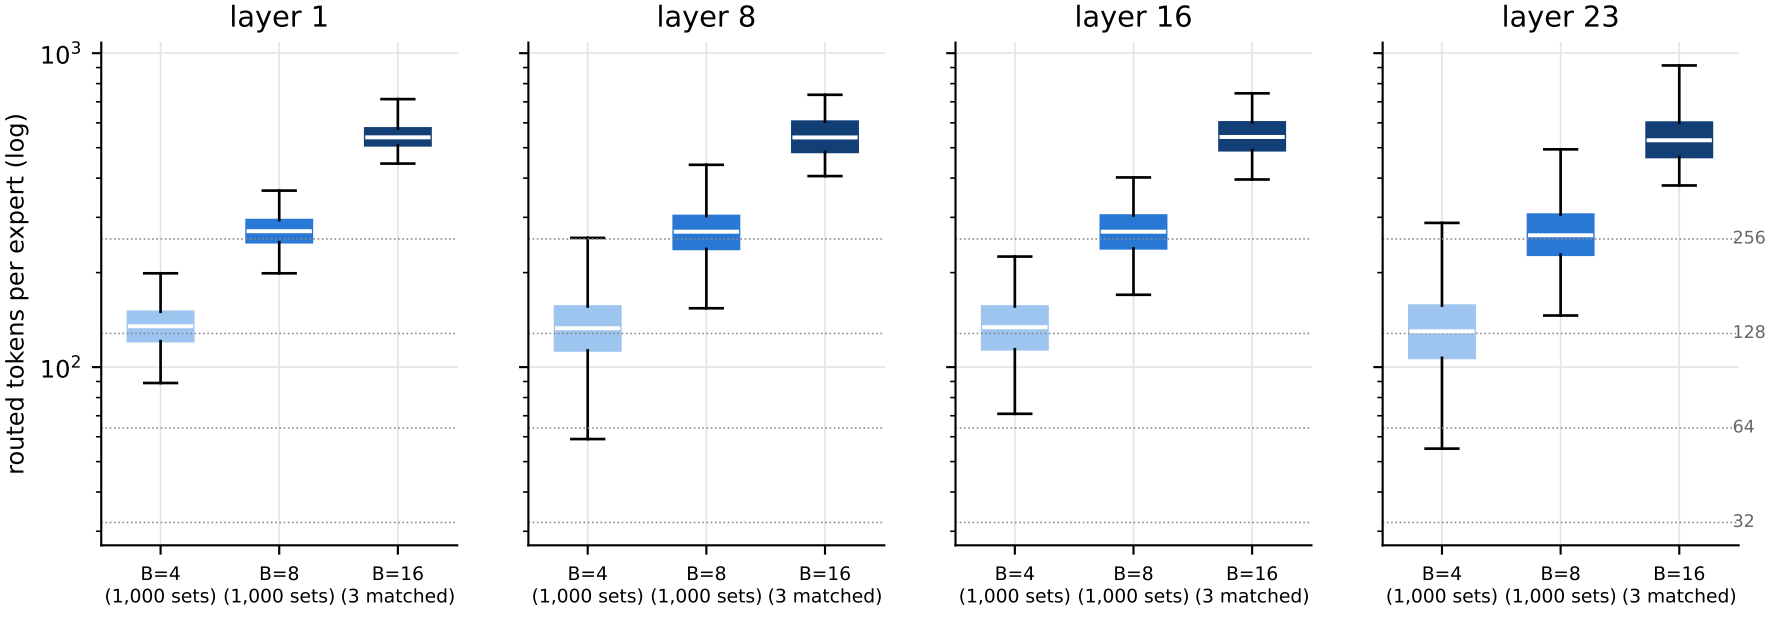
Full Random: per-expert routed-token exposure (whiskers p1–p99; dotted = 32/64/128/256)
layer 1
layer 8
layer 16
layer 23
routed tokens per expert (log)
10^3
10^2
256
128
64
32
B=4
(1,000 sets)
B=8
(1,000 sets)
B=16
(3 matched)